

Enron Submission Questions

Student: Peter Carsten Petersen

P5 - Identify Fraud from Enron Email

1. Summarize for us the goal of this project and how machine learning is useful in trying to accomplish it. As part of your answer, give some background on the dataset and how it can be used to answer the project question. Were there any outliers in the data when you got it, and how did you handle those? [relevant rubric items: "data exploration", "outlier investigation"]

The goal of this project is to build a predictive model, which from a dataset of "Enron" employee's salary items and Email habits can aid in identifying who was a person of interest "poi" in the criminal investigations which followed on the 2002 bankruptcy of the company.

Armed only with this dataset, the task essentially boils down to a highly complex statistical exercise, aimed at identifying possible patterns between the many combinations of features arriving at a choice of algorithm for solving the problem. This could theoretically be done without the assistance of machines but will be a cumbersome and error prone task, and this is where Machine Learning proves to be a powerful tool.

Reaching the project goal employing Machine Learning tools has the following benefits:

1. Calculation time will be barely noticeable and error free compared to a manual process.
2. There are numerous tools available which can aid us in choosing features, scaling data, identifying outliers, splitting datasets for training/testing and reviewing the validity and strength of our result.
3. Armed with this computational power and available tools, we can easily test many possible algorithms, and run many simulations in order to ultimately end at the best possible model for prediction.

Included in the provided dataset is the information if the person is a "poi" or not, i.e. we are already provided with the label of interest. The Project goal is therefore to be reached via a supervised learning algorithm, this will affect the choice of models to be tested.

From visual inspection of the dataset as well as from the exercise in "outliers", we can easily identify and remove 2 particular entries, which are definitely not persons, "Total" which is a summation of all persons data, and "The Travel Agency in the Park". All remaining "outliers" high and low values I have chosen to keep in the dataset. "High" entries since these are most likely to point towards a "POI", and "Low" are equally important to keep as removing them will create a dataset more biased towards only possible "POI's"

2. What features did you end up using in your POI identifier, and what selection process did you use to pick them? Did you have to do any scaling? Why or why not? As part of the assignment, you should attempt to engineer your own feature that does not come ready-made in the dataset – explain what feature you tried to make, and the rationale behind it. (You do not necessarily have to use it in the final analysis, only engineer and test it.) In your feature selection step, if you used an algorithm like a decision tree, please also give the feature importances of the features that you use, and if you used an automated feature selection function like SelectKBest, please report the feature scores and reasons for your choice of parameter values. [relevant rubric items: “create new features”, “properly scale features”, “intelligently select feature”]

Given that the task is to identify “Persons of Interest”, I created and included in the Features_List a calculated value “Bonus Salary Multiple”, i.e. how many times bigger was respective individuals Bonus over Salary, which I suspect would be a strong indicator for “POI” identification, in other words if you were involved in the “Fraud” you would expect to be paid handsomely via excessive Bonus.

From this extended Features_List I have reduced the Features to 10 via “SelectKBest” which resulted in the following list and score per feature.

Feature	Score
Exercised stock options	25,10
Total Stock Value	24,47
Bonus	21,06
Salary	18,58
Deferred Income	11,60
Bonus Salary Multiple (new!)	10,96
Long term incentive	10,07
Restricted Stock	9,35
Total Payments	8,87
Shared receipt with POI	8,75

And finally all chosen features have been scaled to between 0-1 using the SKLearn MINMAX Scaler. While not necessarily needed for all algorithms, it will aid certain Algorithms and not affect other Algorithms, i.e. no worse performance.

3. **What algorithm did you end up using? What other one(s) did you try? How did model performance differ between algorithms? [relevant rubric item: "pick an algorithm"]**

Algorithm	Pre Tuning			Post Tuning and Stratified ShuffleSplit		
	Accuracy Score	Precision Score	Recall Score	Accuracy Score	Precision Score	Recall Score
GaussianNB	0.886	0.5	0.6	0.861	0.444	0.444
Support Vector Classifier (SVC)	0.886	0.0	0.0	0.903	0.750	0.333
RandomForrest	0.841	0.0	0.0	0.868	1.000	0.778

3 different Algorithms were chosen and while the "SVC" and "RandomForrest" performed significantly better on the limited dataset post tuning, they did not survive the tester.py data. The "SVC" on tester.py was abandoned after 18 hours of calculations with no result and "RandomForrest" delivered poor results on tester.py. So in the end the Algorithm which tested well with no "Tuning" and less so with "Tuning" was the only Algorithm which made the necessary scoring in tester.py and is therefore the "clf" reported in this project.

4. **What does it mean to tune the parameters of an algorithm, and what can happen if you don't do this well? How did you tune the parameters of your particular algorithm? (Some algorithms do not have parameters that you need to tune -- if this is the case for the one you picked, identify and briefly explain how you would have done it for the model that was not your final choice or a different model that does utilize parameter tuning, e.g. a decision tree classifier). [relevant rubric item: "tune the algorithm"]**

Parameter Tuning is a method of increasing the accuracy/strength of the Algorithm, if done correctly the ML model will deliver a strong validated product

If done incorrectly it can seriously decrease the speed of using the model and/or can lead to a model with low performance.

For all chosen Algorithms I have applied a GridSearch process, where a range of possible parameters have been tested for all possible combinations in order to find the optimal parameters.

Specifically for the GaussianNB which is the chosen "clf", there are no parameters to tune, instead the strategy here was to do a gridsearch over a range of features (transformed via PCA "Principal Component Analysis") as well as over a further reduction of Features via SelectKBest.

5. What is validation, and what's a classic mistake you can make if you do it wrong? How did you validate your analysis? [relevant rubric item: "validation strategy"]

Validation is the process of assessing the strength of the Machine Learning Model. This is done via splitting the data into training sets and test sets, so that the Model can be "validated" on the test set after being fitted on the training set. Among classic mistakes can be testing on same data as model has been trained, allowing too many features/classifications of data "Overfitting" the model, or only validating on the "Accuracy" score.

Ideally validation should always be done via splitting data in training/test sets, this done practically by allocating unique portions of data to training and test sets respectively, or running many fitting/validating exercises on Model via randomly splitting data many times in training/test sets. The latter option has been used for this exercise via a "StratifiedShuffleSplit" when running chosen ML algorithms with parameter tuning.

6. Give at least 2 evaluation metrics and your average performance for each of them. Explain an interpretation of your metrics that says something human-understandable about your algorithm's performance. [relevant rubric item: "usage of evaluation metrics"]

Algorithm	Pre Tuning			Post Tuning and Stratified ShuffleSplit			Post Tuning and validation in tester.py		
	Accuracy Score	Precision Score	Recall Score	Accuracy Score	Precision Score	Recall Score	Accuracy Score	Precision Score	Recall Score
Gaussian NB	0.886	0.5	0.6	0.861	0.444	0.444	0.843	0.394	0.334

Precision Score: Tells us how many of the "POI's" the ML model has predicted in the test set, which were actually classified as "POI's". Score is between 0 (no "POI's" successfully classified) to 1 (all "POI's" predicted were actual "POI's"). This score does not tell us how many of the "POI's" were found.

Recall Score: Tells us how many of the "POI's" the ML model has predicted in the test set, vs. the Actual amount of "POI's" in the test set. Score is between 0 (no "POI's" in the test set successfully classified) to 1 (all "POI's" in the test set successfully classified). This score does not tell us if all predicted "POI's" were actual "POI's".

So in order to conclude if ML Model is strong it takes a combination of scores close to 1 for both Precision and Recall. I.e. amount of "POI's" predicted had a high Precision and amount of Actual "POI's" predicted was a larger portion of total - Recall.

References:

SKLearn documentation

Udacity Forum discussions

Github repositories on Machine Learning

Wikipedia